

Training pool vs CRUXEval-O test set — trace-complexity distributions

Bars = within-set share; each bar also labelled with its absolute count (n). Bins stay fine and the tail is kept segmented to each set's maximum — note the test set's trace-length tail out to ≈28k tokens that the ≤3072-filtered training pool lacks.

